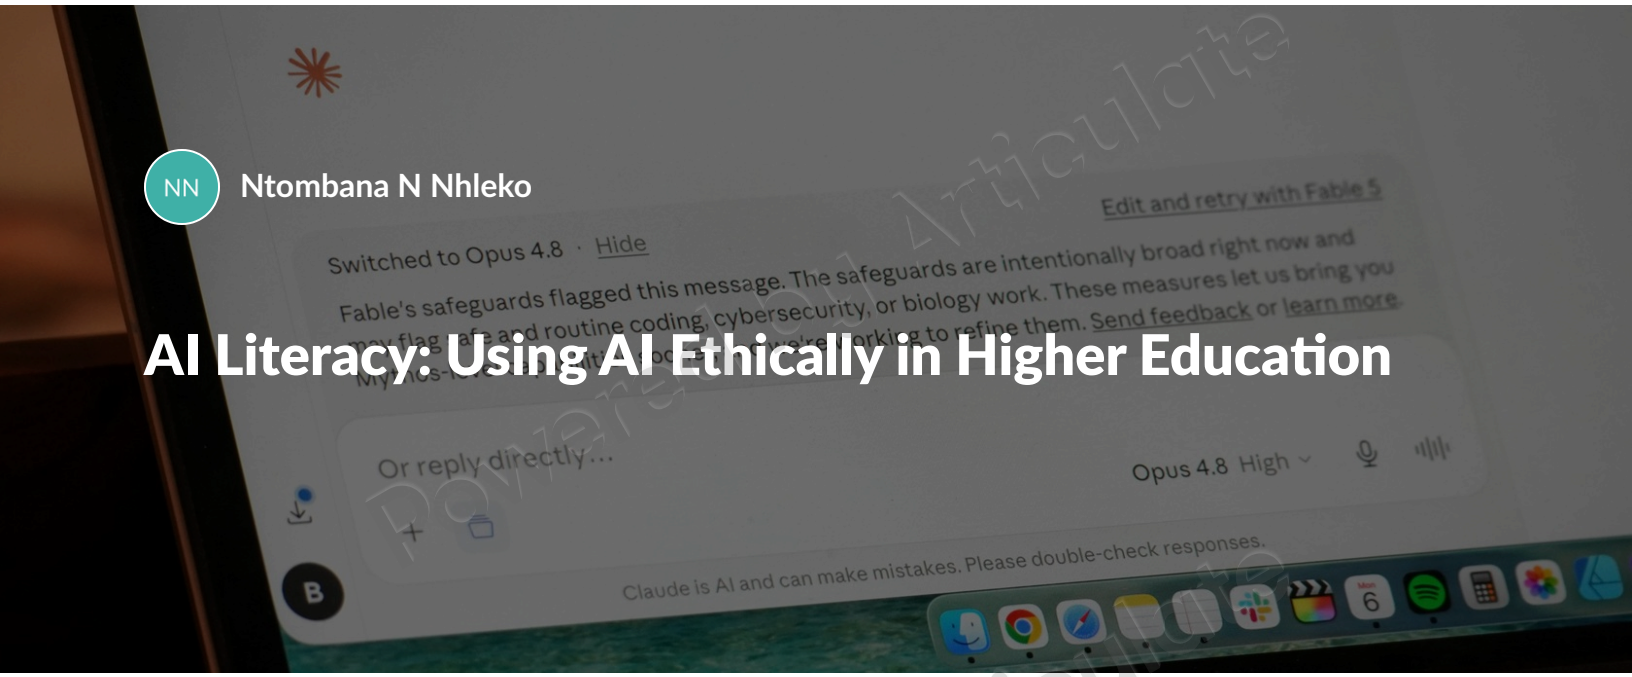

AI Literacy: Using AI Ethically in Higher Education

Powered by Articulate

Course Description

This course is a practical, five-module course designed to help students understand generative AI and use it responsibly in their academic work. Students will explore AI tools for learning and research, identify risks such as misinformation, bias, plagiarism and fabricated citations, understand institutional AI-use guidelines, and apply strategies for using AI to support learning without replacing their own critical thinking, academic integrity or intellectual contribution.

Learning outcomes

By the end of this course, students will be able to:

1. **Identify and selec**t appropriate AI tools for academic learning, research and writing tasks.
2. **Recognise and evaluate** risks associated with AI-generated content, including misinformation, bias, plagiarism and false citations.
3. **Apply** institutional guidelines for ethical AI use, including transparency, disclosure and academic integrity requirements.

4. **Use AI responsibly** to support brainstorming, understanding concepts, assessment feedback and independent study.
5. **Evaluate their use of AI** and make informed decisions that preserve their own critical thinking, academic judgement and ownership of their work.

- The New Landscape
- Ethical Considerations
- Quiz 1
- University Policy and Guidelines on AI Use
- Quiz 2
- Using AI to Support Learning
- Your AI Decision
- Student Survey

Powered by Articulate
Powered by Articulate
Powered by Articulate

Lesson 1 of 8

The New Landscape

NN Ntombana N Nhleko

Powered by Articulate

This lesson introduces generative AI: what it is, its origins and the growing ecosystem of tools built on it. By the end of this lesson students will be able to identify the different types of tools and determine the right category of tool for a given academic task, and understand how to prompt these tools responsibly.

Introduction to Generative AI

WHAT IS GENERATIVE AI | A BRIEF HISTORY OF AI | AI TOOLS FOR LEARNING, RESEARCH AND WRITING

Generative AI refers to AI systems that can create new content — such as text, images, audio, or video — based on the inputs (prompts) they are given. Rather than simply retrieving existing information, generative AI models produce new output by recognising and reproducing statistical patterns learned from very large datasets.

Commonly used examples of generative AI include chatbots such as ChatGPT and Claude, voice assistants such as Siri or Google Assistant, and tools that generate detailed text, images, or code

from a prompt.

Understanding where generative AI came from helps put today's tools in context. AI is not a new idea — it is the product of decades of incremental research, punctuated by a few major breakthroughs.

Take a look at this infographic by Tech Target Network to see the evolution and progression of AI development:

Powered by Articulate
Powered by Articulate
Powered by Articulate

Generative AI's evolution

For an advanced technology that's considered relatively new, generative AI is deep-rooted in history and innovation.

1932
Georges Artsrouni invents a machine he reportedly called the "**mechanical brain**" to translate between languages on a mechanical computer encoded onto punch cards.

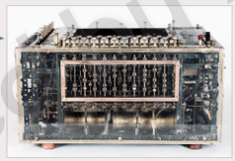

1957
Linguist **Noam Chomsky** publishes *Syntactic Structures*, which describes grammatical rules for parsing and generating natural language sentences.

1966
MIT professor Joseph Weizenbaum creates the first chatbot, **Eliza**, which simulates conversations with a psychotherapist.

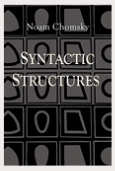

1968
Computer science professor Terry Winograd creates SHRDLU, the first multimodal AI that can manipulate and reason out a world of blocks according to instructions from a user.

1980
Michael Toy and Glenn Wichman develop the Unix-based game *Rogue*, which uses procedural content generation to dynamically generate new game levels.

1985
Computer scientist and philosopher Judea Pearl introduces Bayesian networks causal analysis, which provides statistical techniques for representing uncertainty that leads to methods for generating content in a specific style, tone or length.

1986
Michael Irwin Jordan lays the foundation for the modern use of recurrent neural networks (RNNs) with the publication of "Serial order: a parallel distributed processing approach."

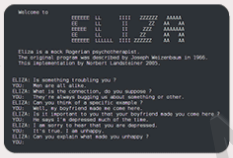

1989
Yann LeCun, Yoshua Bengio and Patrick Haffner demonstrate how convolutional neural networks (CNNs) can be used to recognize images.

2000
University of Montreal researchers publish "A Neural Probabilistic Language Model," which suggests a method to model language using feed-forward neural networks.

2006
Data scientist Fei-Fei Li sets up the ImageNet database, which provides the foundation for visual object recognition.

2011
Apple releases **Siri**, a voice-powered personal assistant that can generate responses and take actions in response to voice requests.

2012
Alex Krizhevsky designs the AlexNet CNN architecture, pioneering a new way of automatically training neural networks that take advantage of recent GPU advances.

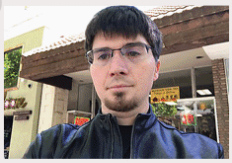

2013
Google researcher Tomas Mikolov and colleagues introduce word2vec to identify semantic relationships between words automatically.

2014
Research scientist **Ian Goodfellow** develops generative adversarial networks (GANs), which pit two neural networks against each other to generate increasingly realistic content.

Diederik Kingma and Max Welling introduce variational autoencoders to generate images, videos and text.

2015
Stanford researchers publish work on **diffusion models** in the paper "Deep Unsupervised Learning using Nonequilibrium Thermodynamics." The technique provides a way to reverse-engineer the process of adding noise to a final image.

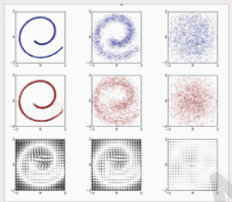

2017
Google researchers develop the concept of transformers in the seminal paper "Attention is all you need," inspiring subsequent research into tools that could automatically parse unlabeled text into large language models (LLMs).

2018
Google researchers implement transformers into BERT, which is trained on more than 3.3 billion words and can automatically learn the relationship between words in sentences, paragraphs and even books to predict the meaning of text. It has 110 million parameters.

Google DeepMind researchers develop AlphaFold for predicting protein structures, laying the foundation for generative AI applications in medical research, drug development and chemistry.

OpenAI releases GPT (Generative Pre-trained Transformer). Trained on about 40 gigabytes of data and consisting of 117 million parameters, GPT paves the way for subsequent LLMs in content generation, chatbots and language translation.

2021
OpenAI introduces **Dall-E**, which can generate images from text prompts. The name is a combination of WALL-E, the name of a fictional robot, and the artist Salvador Dali.

2022
OpenAI releases **ChatGPT** in November to provide a chat-based interface to its GPT 3.5 LLM. It attracts over 100 million users within two months, representing the fastest ever consumer adoption of a service.

2023
The release of **GPT-4** sparked the multimodal revolution, showcasing impressive capabilities in processing both text and visual data.

2023
Microsoft integrates a version of ChatGPT into its Bing search engine. Google releases the Bard chat service based on its Lamda engine. The controversy over detecting AI-generated content continues.

2024
Multimodal AI models gain traction. Agentic AI goes mainstream, spurred by Salesforce's announcement of a major pivot

2025
Chinese AI startup **DeepSeek** claims its new

Generative AI is not a single tool, it is a growing/evolving category made up of tools with very different purposes. Knowing what each tool is designed for (and, just as importantly, what it is not designed for) is the first step to using AI responsibly in your academic work.

Tool	Category	What it does	Where it fits in the learning process
ChatGPT	Generative AI chatbot	A general-purpose conversational AI that can brainstorm, explain concepts, summarise, draft, and answer questions across almost any topic.	Brainstorming and idea generation; clarifying concepts; outlining. Not a source of citable academic evidence, and should never generate your final submitted argument or text.
Perplexity	AI-powered search / answer engine	Combines a language model with live web search, generating an answer alongside links to the sources it drew on.	Early-stage research and orientation on a topic. Useful for identifying possible sources to verify independently — the underlying sources must still be checked for credibility, and not taken at face value.

Powered by Articulate

Grammarly	Editorial / writing-assistance software	Checks spelling, grammar, and punctuation; suggests alternative word choices; can comment on clarity, tone, and sentence structure.	Proofreading and polishing your own writing — similar to a spellchecker. Should support, not replace, your own writing voice and argument.
ResearchRabbit	Reference / literature-mapping tool	A free "citation graph" tool that visually maps how academic papers relate and cite one another, helping you discover related literature from a starting paper.	Literature searching and discovery — finding relevant, credible peer-reviewed sources once you already have a seed article or topic.
Trinka	AI-powered academic editing tool	An editor built specifically for academic and scientific writing, checking grammar, tone, and adherence to academic style conventions (including discipline-specific language checks).	Refining the language and academic tone of a draft you have already written — not for generating the content itself.
Semantic Scholar	Reference management / academic search engine	An AI-powered academic search engine that indexes millions of papers, extracts key findings, and helps you filter for peer-reviewed, credible literature.	Finding, organising, and evaluating academic sources for a research task or literature review.

Powered by Articulate
Powered by Articulate
Powered by Articulate

ReadCube (Papers)	Reference management software	Helps you organise, annotate, and manage the academic PDFs and citations you collect during research, and can generate correctly formatted reference lists.	Managing your reference library and formatting citations once you have found credible sources — it does not generate or verify the content of sources for you.

Powered by Articulate

There is a pattern across this table: editorial tools (Grammarly, Trinka) support the language of work you have already written; reference tools (ResearchRabbit, Semantic Scholar, ReadCube) support finding and organising credible sources and generative chatbots (ChatGPT, Perplexity) support thinking, brainstorming, structuring and explaining — but should never be the source of your final academic argument or citations.

Powered by Articulate

Prompting questions - what to ask and what not to ask.

Powered by Articulate

Questions that are appropriate to ask AI tools

- **"Can you suggest three possible approaches I could take on this essay topic?"**
- **"Can you explain this concept in simpler terms, using an example?"**
- **"Can you check this paragraph for grammar and clarity?"**

- **"What are some counterarguments to this position that I should consider?"**
- **"Can you help me understand why my in-text citation formatting is incorrect?"**
- **"Quiz me on this topic using short-answer questions."**

Questions and requests to avoid

- **"Write my essay/assignment for me."**
- **"Give me a full paragraph I can copy into my assignment."**
- **"Find me academic sources on this topic" (without independently verifying that each source actually exists and is credible).**
- **"Rewrite this so it won't be flagged by Turnitin/AI detection."**
- **"Generate a reference list for this essay" (AI can fabricate convincing-looking but non-existent references — see lesson 2).**
- **Sharing another student's, client's, or research participant's private or identifiable information in a prompt.**

To determine whether your prompt is good or not, ask yourself: if the response you are hoping for is the actual words, argument, or citations that will appear in your submitted work, you are asking the AI to do your work for you — not to assist your thinking.

Powered by Articulate

Knowledge Check

Sort the various actions into two categories - **Acceptable AI Use** or **Academic Misconduct**

Acceptable AI Use

Using Generative AI to brainstorm an essay outline or check for grammatical flow

Using Grammarly to correct sentence structure and syntax in an original paper

Academic Misconduct

Powered by Articulate
Powered by Articulate
Powered by Articulate

Using AI to source academic sources and articles without verifying them

Copying an AI-generated paragraph directly into a final draft without citations

Powered by Articulate
Powered by Articulate
Powered by Articulate

Ethical Considerations

NN Ntombana N Nhleko

Powered by Articulate

Generative AI tools are powerful, but they are not neutral or inherently honest and reliable. This lesson covers three of the most important ethical risk areas when using AI in academic work: misinformation and bias, plagiarism, and false citations. Generative AI tools can falsely generate information which does not exist/is not factually correct or biased information. It is important for students to understand how to verify accuracy of information.

Powered by Articulate

Powered by Articulate

Misinformation and Bias

Generative AI models are trained on massive datasets drawn largely from the internet. This means the model's "knowledge" reflects whatever was in that data, including its errors, gaps, one-sided perspectives, and prejudices.

How can AI push misinformation and bias into its responses?

- **Confident-sounding errors/ "Hallucinations":** AI models are designed to generate plausible, fluent text — not necessarily true text. They will often state incorrect facts, outdated information or invented details with the same confident tone as accurate information.
- **Training-data bias:** If certain viewpoints, demographics, or cultural perspectives are over- or under-represented in the training data, the model's output can unintentionally favour some perspectives over others, or reflect stereotypes present in that data presented as fact.
- **Outdated knowledge:** Most models have a training "cut-off" date and may not reflect recent developments, corrections, or retracted research unless they have live web access.
- **Lack of source transparency:** Many generative AI tools cannot reliably tell you where a specific claim came from, making it hard to trace and verify the underlying evidence.
- **Reinforcement through repetition:** Because millions of people use the same underlying models, a biased or inaccurate framing generated by AI can be repeated and normalised at scale if it isn't critically checked.

Powered by Articulate

Example

> A student asks an AI chatbot to summarise "the causes of a historical event." The response is fluent and confident, but subtly overrepresents one country's perspective because that viewpoint was more heavily represented in the training data, and omits several sources critical of that dominant narrative. Without independently checking peer-reviewed historical sources, the student unknowingly reproduces a one-sided account in their assignment.

Practical safeguard: always cross-verify AI-generated information against reliable, peer-reviewed academic sources and your own research. Actively ask yourself whose perspective might be missing from the response you were given.

Plagiarism

Plagiarism in AI Tools

- Submitting AI-generated text as your own original work, without disclosure, is a form of plagiarism — even though there is no single named "author" being copied from.
- AI models are frequently trained on copyrighted and previously published material (including books, journal articles, and websites) without always clearly indicating this. Content generated by AI can therefore closely mirror existing copyrighted text without you realising it.
- Ghostwriting: relying solely on AI to generate the argument, analysis, or structure of an assessment — even if you then lightly

Powered by Articulate

edit it — constitutes ghostwriting and breaches academic integrity, because the core intellectual work was not yours.

- Self-plagiarism risk: using AI to "reword" your own previously submitted work can also raise self-plagiarism concerns if it is resubmitted as new work without appropriate disclosure.
- Translation plagiarism: using AI to translate someone else's published work and presenting the translation as original writing is still plagiarism.

Most Institutions increasingly require transparency rather than an outright ban: if you use AI as part of your process, you are generally expected to disclose this (see Lesson 3), not to conceal it.

False Citations

- AI language models can generate references that look completely legitimate with plausible author names, journal titles, years, and even DOIs — that do not actually exist. This is sometimes called "citation hallucination."
- AI can also misattribute real findings to the wrong author, journal, or year, or blend details from two different real sources into one fabricated reference.
- Because AI-generated citations look correctly formatted, they can pass a casual glance and only be caught through independent verification.

Powered by Articulate
Powered by Articulate
Powered by Articulate

- Submitting fabricated or unverified citations is treated as a serious academic integrity violation, regardless of whether the student intended to deceive — the responsibility to verify sources rests with the student.

Golden rule

- Never cite a source you have not personally located, opened, and read. If AI suggests a source, treat it only as a lead — verify it exists (via a database like Semantic Scholar, Google Scholar, your digital library or a simple web search) and also confirm the source actually says what the AI claims before citing it.
- Do not cite ChatGPT, or similar generative AI tools, as an academic or evidentiary source — most institutions and citation bodies (including APA) do not consider AI chat output a credible academic reference, even where preliminary guidance exists for how to cite an AI tool as the origin of specific generated text.

Powered by Articulate

Lesson 3 of 8

Quiz 1

NN Ntombana N Nhleko

Test your understanding of lesson 2 by completing this 3-item quiz.

Powered by Articulate
Powered by Articulate
Powered by Articulate

Question

01/03

You use Perplexity AI for an assignment. The AI tool generates information and references - this is one of the sentences provided by the AI tool in the generated response/content:

“Individuals diagnosed with bipolar disorder must always be hospitalised at a mental health institution for at least 30 days before being prescribed medication”.

Is this acceptable

---

1. Yes - Perplexity AI said so therefore it must be accurate and can go straight into the assignment.

2. Yes - as long as I paraphrase the sentence in my own words.

3. No, the claim is inaccurate and reinforces harmful stereotypes; it must be checked against credible peer-reviewed sources before use.

4. No, but only because Perplexity is a search style tool - a chatbot like ChatGPT would be trustworthy here.

Powered by Articulate

Question

02/03

You copy and paste AI-generated paragraphs directly into your assignment and submit them along with some of your original work. You have also disclosed the prompting information/response from the AI tool.

Is this acceptable?

---

1. Yes — because I disclosed the AI tool and included the prompt sheet, this fully complies with academic integrity policy.

2. No — copying and pasting AI-generated text directly into an assessment is plagiarism/ghostwriting even when disclosed; the wording and ideas submitted must be your own.

3. Yes — since only part of the assignment was AI-generated and the rest was my own writing.

4. No, but only because I didn't cite ChatGPT as a source.

Powered by Articulate
Powered by Articulate
Powered by Articulate

Question

03/03

While doing research for an assignment, you use an AI tool which suggests several academic sources. When you try to locate them in the library database or online, three of the five sources cannot be found anywhere.

What must you do next?

---

1. Trust them anyway — AI tools verify facts before responding, so the sources must exist somewhere.

2. Use them, but adjust the publication year so they look more current.

3. Do not use them — AI can generate convincing but non-existent ("hallucinated") citations; every source must be independently located, opened, and verified before it's cited.

4. It's fine to cite them as long as the reference is formatted correctly in APA style.

Powered by Articulate

University Policy and Guidelines on AI Use

NN Ntombana N Nhleko

Powered by Articulate

Every institution's specific rules may differ, but most higher-education AI policies — including the policy summarised below — share a common underlying philosophy: AI should enhance learning, not replace it, and its use must be transparent.

University Policy on AI Use

The institution's overall position

- Artificial Intelligence is recognised as a significant and transformative technology for learning, teaching, assessment, and research — not something to be ignored or blanket-banned.
- Students should use AI to ENHANCE their personal learning, not replace it.
- AI-generated content must always be critically evaluated for accuracy and potential bias, and cross-checked against reliable academic sources.

- AI must never be the sole source of information for, or the author of academic work.

Academic integrity and Plagiarism

- Presenting AI-generated content as one's own original work, without disclosure, is treated as a breach of academic integrity.
- Relying solely on AI to generate an argument for an assessment constitutes ghostwriting.
- Students remain fully responsible for verifying the academic credibility and accuracy of any source or data identified with AI assistance.

Transparency and Disclosure requirements

- Students must clearly acknowledge when and how AI tools were used, and be transparent about which parts of the work are their own contribution.
- Students are required to submit an AI prompt sheet alongside relevant assessments, documenting: the date and time of use, which AI tool was used, the exact prompts given, and the responses received.
- Where AI is used as an editing tool, before-and-after (pre- and post-) screenshots will be required to show what was changed.
- AI and grammar/paraphrasing tools (e.g., Grammarly, QuillBot, ChatGPT) must be disclosed and where applicable, included in the reference list.

Powered by Articulate

Citation Rules

- Quoting or citing an AI chatbot (e.g., ChatGPT) as an academic source or piece of evidence is not accepted practice.
- All information sourced with AI assistance must be independently cross-checked against peer-reviewed, published academic literature, which must then be cited instead.
- Fabricated or unverifiable citations — whether AI-generated or not — are treated as an academic integrity violation.

Privacy, Bias and Fairness

- Students should be mindful of privacy and data security when using AI tools — avoid entering confidential, personal, or identifiable information (your own or others') into AI systems.
- Students are expected to critically evaluate AI output for potential bias and to strive to produce fair, inclusive work, rather than uncritically accepting AI-generated framing.

Powered by Articulate
Powered by Articulate
Powered by Articulate

What to do if you are unsure?

- Consult your institution's academic integrity policy and student handbook for the full, current wording —note that policies are updated as AI technology evolves.

- Speak to an academic advisor or lecturer before using AI in a way you are unsure about — asking first is always safer than assuming.
- If you are uncomfortable using generative AI for personal or ethical reasons, you are not required to use it. It should always be treated as an optional aid to learning, not a mandatory replacement for your own effort.

Powered by Articulate

Practical checklist before submitting

- [ ] Have I disclosed every AI tool I used, and for what purpose?
- [ ] Have I completed and attached an AI prompt sheet, where required?
- [ ] Have I personally verified every citation and source in my reference list?
- [ ] Could I explain and defend my central argument, unaided, if required?
- [ ] Have I checked my draft through the institution's similarity-checking tool before final submission?

Powered by Articulate

Powered by Articulate

Lesson 5 of 8

Quiz 2

NN Ntombana N Nhleko

Test your understanding of lesson 3 by completing this 3-item quiz.

Powered by Articulate
Powered by Articulate
Powered by Articulate

Question

01/03

According to university policy, which of the following counts as a breach of academic integrity?

---

1. Presenting AI-generated content as your own original work without disclosing it.
2. Submitting a draft through the Turnitin draft checker before final submission.
3. Using AI to check grammar and then disclosing this in the reference list.
4. Cross-checking an AI-suggested source against peer-reviewed literature.

Powered by Articulate
Powered by Articulate
Powered by Articulate

Question

02/03

What must students do to remain transparent about AI use in an assignment?

- 1. Only mention AI use if directly asked by the marker.
- 2. Cite the AI chatbot itself as an academic source in the reference list.
- 3. Submit an AI prompt sheet documenting the tool used, exact prompts, and responses received, plus before-and-after screenshots if AI was used for editing.
- 4. Nothing, as long as the final work reads well.

Powered by Articulate
Powered by Articulate
Powered by Articulate

Question

03/03

Under the citation rules in the policy, what should a student do if AI helps them find a useful fact or claim?

---

1. Cite the AI chatbot directly as the source of the fact.

2. Include the fact without any citation, since AI already checked it.

3. Assume the fact is accurate as long as it sounds plausible.

4. Independently verify the fact against peer-reviewed, published academic literature and cite that instead.

Powered by Articulate
Powered by Articulate
Powered by Articulate

Using AI to Support Learning

NN Ntombana N Nhleko

Powered by Articulate

When used well, generative AI can be a genuinely valuable study partner. This lesson focuses on the constructive and ethical ways AI can support four common learning activities: brainstorming, explaining concepts, giving feedback on assessments and general study support. Understanding how to utilise AI tools in this digital age is a necessary skill to acquire. Even if one does not intend to utilise it, it is beneficial to know how AI tools can typically aid in academic study.

AI Contribution Continuum

So, where does this all fit into the learning process? A useful way to think about this is a continuum of AI contribution versus your own contribution/original thinking to a piece of work - the AI literacy continuum is a framework for understanding how university students engage with generative AI across different academic tasks. The continuum is based on the idea that: the value of AI use cannot be judged only by how much AI a student uses. Instead it depends on the balance between independent human thinking and purposeful, critically evaluated AI support. A student who uses AI extensively may still be learning effectively if they remain the author of their ideas, verify outputs and use AI to deepen their understanding. Conversely,

a student who submits AI-generated work with little engagement may appear productive while actually avoiding the intellectual work that university higher education study is intended to develop.

The continuum is also informed by research on productive failure and desirable difficulty, which suggests that effort, uncertainty, error and revision are often necessary for deep learning. When learners first attempt a problem, from an argument, or explain a concept before receiving support, they are more likely to recognise gaps in their understanding and develop transferable knowledge. Generative AI can reduce this productive struggle if it is used too early or too heavily to produce finished answers. This may lead to what Fan et al. (2025) describe as “metacognitive laziness”. Where task performance improves but students monitor, evaluate and direct their own learning less effectively. For this reason, the continuum does not treat AI as inherently good or bad. Its educational value depends on the task, the student’s current level of knowledge, the assessment requirements and whether AI is being used to support thinking or replace it.

Powered by Articulate
Powered by Articulate
Powered by Articulate

The Continuum

The continuum moves from low engagement and low independent thinking towards deliberate, critical and ethical AI use. It should be read as a description of a student's approach to a particular task, rather than a permanent label. For example, a student may demonstrate strong critical AI use when preparing an essay, but rely uncritically on AI when learning a new statistical method.

Powered by Articulate

Neglected [Low human engagement, low AI use]

The student neither engages independently with the task nor uses the AI tool purposefully. Work may be superficial, incomplete or produced without understanding. This stage indicates a need for academic support, clarity about task expectations or improved access to and use of appropriate tools.

Dependent [Low human engagement, high AI use]

The student asks AI to generate most of the content and submits or lightly edits the output. There is limited independent reasoning, verification or ability to explain the work. AI has replaced, rather than supported, student's learning.

Uncritical [Some human engagement, high AI use]

The student uses AI substantially but accepts its outputs with limited scrutiny. They may reword, format or grammatically edit text but do not consistently check accuracy, sources, bias, relevance or disciplinary appropriateness.

Supported Learning [High human engagement, some AI use]

The student begins with their own thinking - such as brainstorming, planning, drafting, solving or forming a position and then uses AI selectively to test ideas, explore counterarguments, clarify concepts, receive feedback or identify possible sources. The student remains the primary author and decision maker.

Critical Partnership [High human engagement, purposeful AI use]

The student deliberately decides when AI is helpful and when independent effort is necessary. They evaluate AI outputs against credible sources and disciplinary standards, identify errors, bias and limitations, disclose use appropriately and can explain and defend their final work. At this stage students may also design effective prompts or AI-supported workflows.

The Continuum in Practice

The appropriate stage will vary according to the task. For example, during the early stages of writing an essay, a student should normally begin in **Stage 3** by developing their own argument or outline before asking AI for feedback or alternative perspectives. When checking references, understanding a difficult concept, or improving clarity, AI may play a larger role provided the student verifies the information and remains responsible for the final work. In contrast, using AI to generate a complete assessment response would usually fall into **Stage 1**, because it removes the student's opportunity to practise analysis, argumentation and academic writing.The goal of this course is therefore not to encourage students to avoid AI, nor to use it as much as possible. The goal is to help students move towards **Stage 3 and Stage 4**: using AI as a transparent, ethical, and critical learning partner while retaining ownership of their ideas, learning, and academic judgement. Students should always begin by checking the AI-use rules for their module and assessment, because permissible AI use may differ across courses, tasks, and disciplines.

How to: Support Learning using AI Tools

Watch this video which covers the ways in which AI can be used for learning.

Powered by Articulate

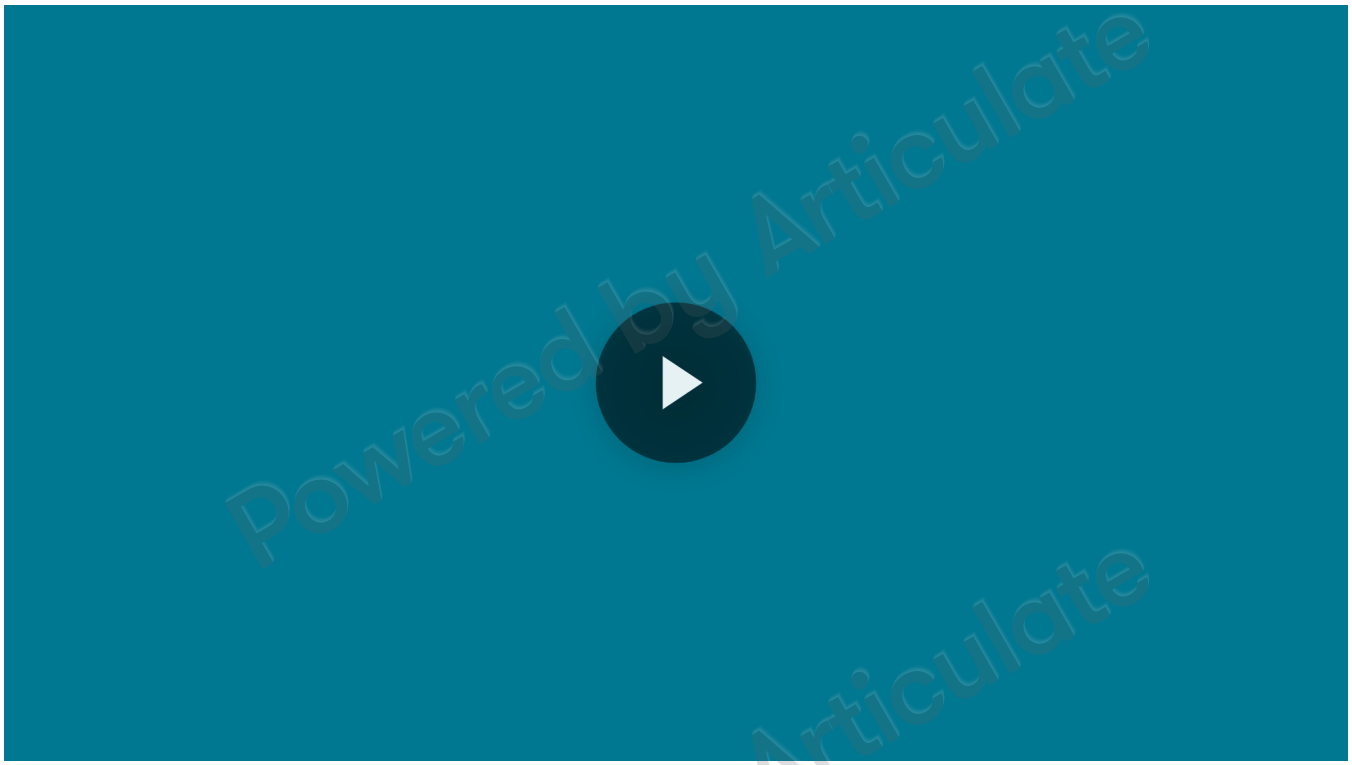

Worksheet

Download and return back to this worksheet when making use of AI tools for academic related tasks.

DOC

Put it into practice Using AI to Support Learning.docx

1.1 MB

Powered by Articulate

Lesson 7 of 8

Your AI Decision

NN Ntombana N Nhleko

Powered by Articulate

"As a student, you are actively practising to be the type of professional you wish to become." — **Michelle Obama**

Across all lessons, one principle holds constant: generative AI is a powerful aid to thinking, research, and productivity — but it is not a substitute for your own critical thinking, original voice, or academic integrity. Used transparently and critically, AI can enhance your learning journey. Used uncritically or dishonestly, it undermines the very skills your education is meant to build.

This capstone assessment will review the concepts covered in the 5 lessons and measure your understanding.

Powered by Articulate

Powered by Articulate

Matching Activity

Match each AI tool to the academic task it is best suited for (refer back to lesson 1).

ResearchRabbit	Brainstorming possible angles or outlines for an assignment.
ReadCube	Early-stage research. Generating an answer with linked sources to verify independently.
ChatGPT	Proofreading your own writing for grammar, punctuation, and tone.
Grammarly	Refining the academic tone and discipline-specific style of a draft you have already written.
Trinka	Mapping how academic papers cite one another, starting from a seed article.
Perplexity	Searching and filtering millions of papers for peer-reviewed and credible literature

Powered by Articulate

Semantic Scholar	Organising your reference library and generating formatted citation lists.

SUBMIT

Scenario

Read each scenario carefully and use the AI continuum framework (covered in Lesson 4) to choose the most appropriate decision for the student to make in each scenario.

Powered by Articulate
Powered by Articulate
Powered by Articulate

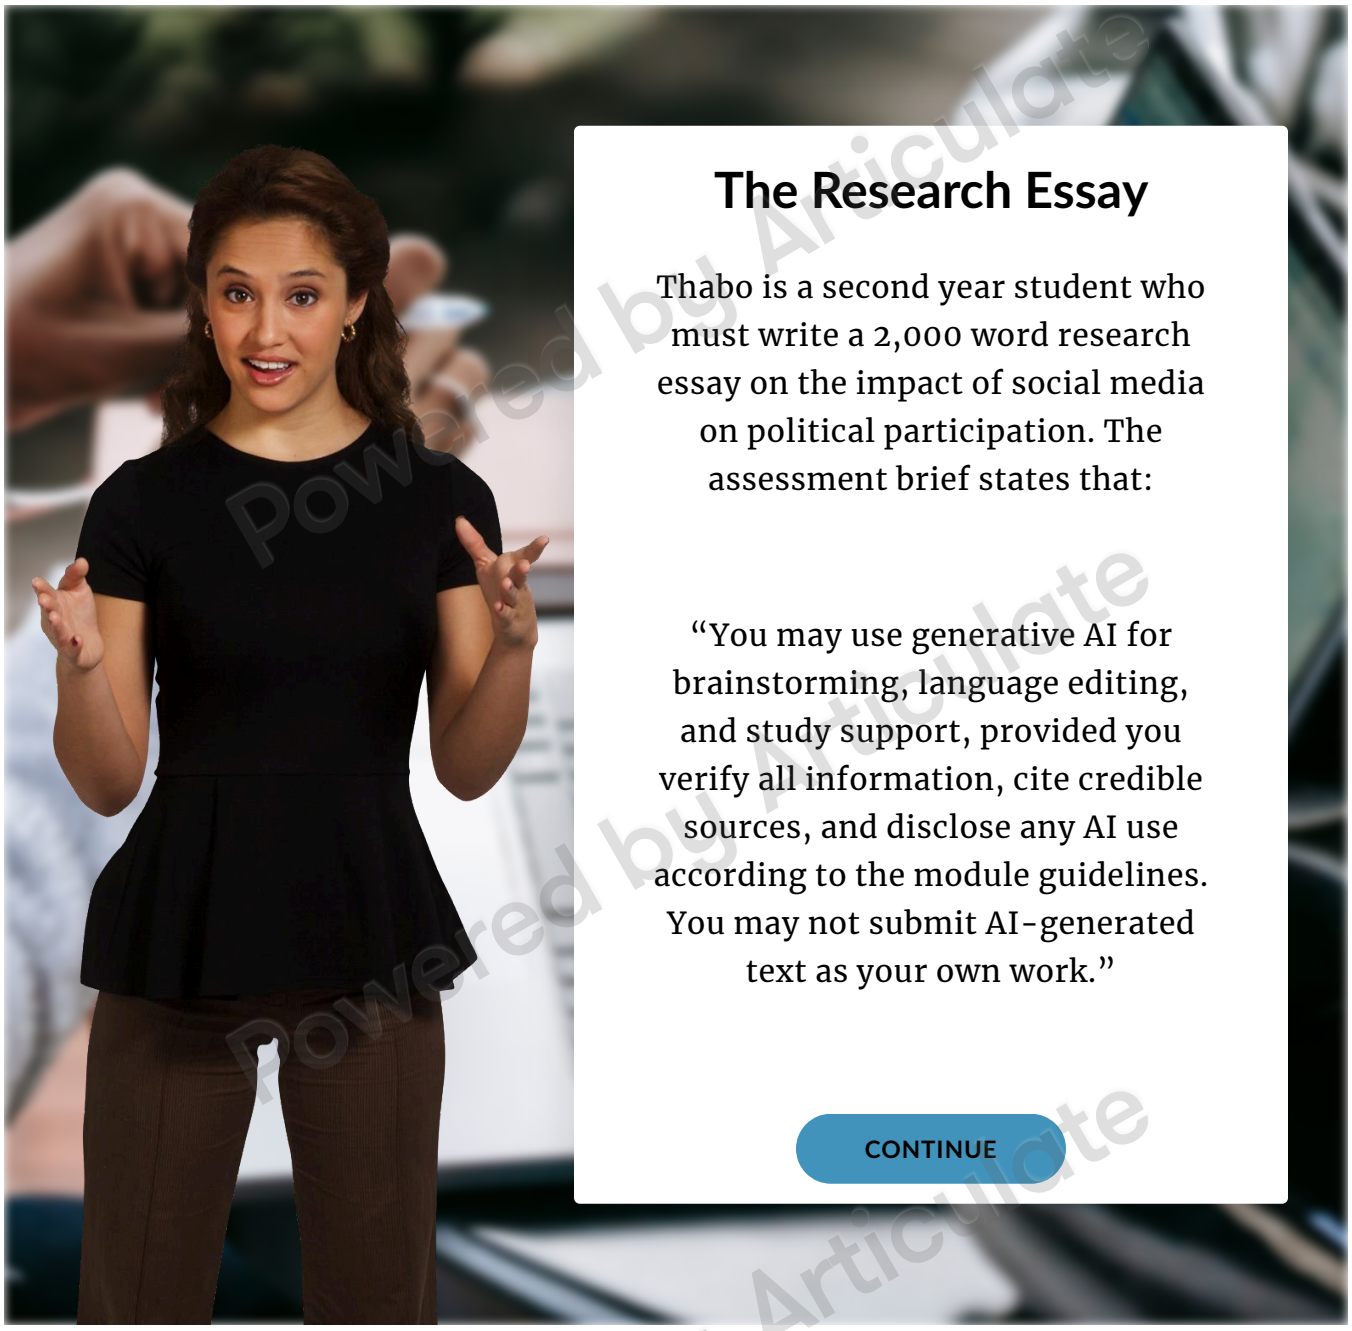

Powered by Articulate
Powered by Articulate
Powered by Articulate

Scene 1 Slide 1

Continue → Next Slide

Powered by Articulate

Scene 1 Slide 2

0 → Next Slide

1 → Next Slide

2 → Next Slide

Scene 1 Slide 3

0 → Next Slide

1 → Scene 1 Slide 1

2 → Next Slide

Powered by Articulate

Scene 1 Slide 4

0 → Next Slide

1 → Scene 2 Slide 1

2 → Next Slide

Powered by Articulate

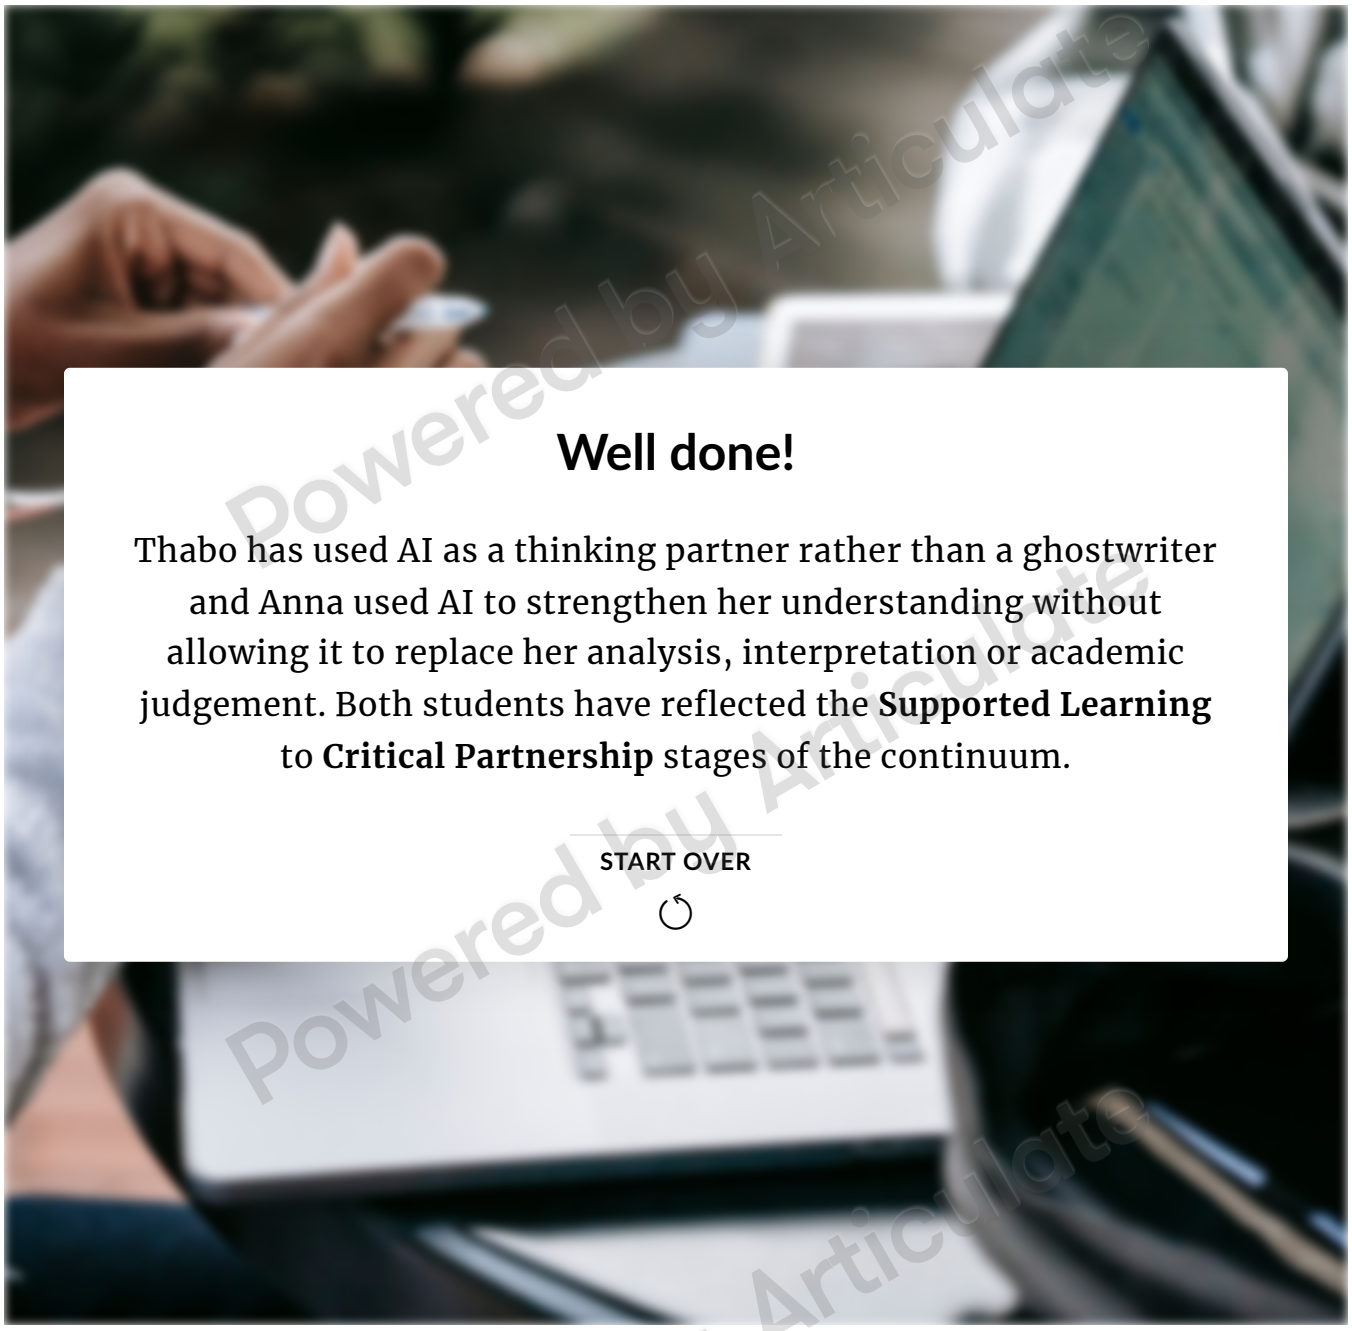

Powered by Articulate
Powered by Articulate
Powered by Articulate

Scene 1 Slide 5

Continue → End of Scenario

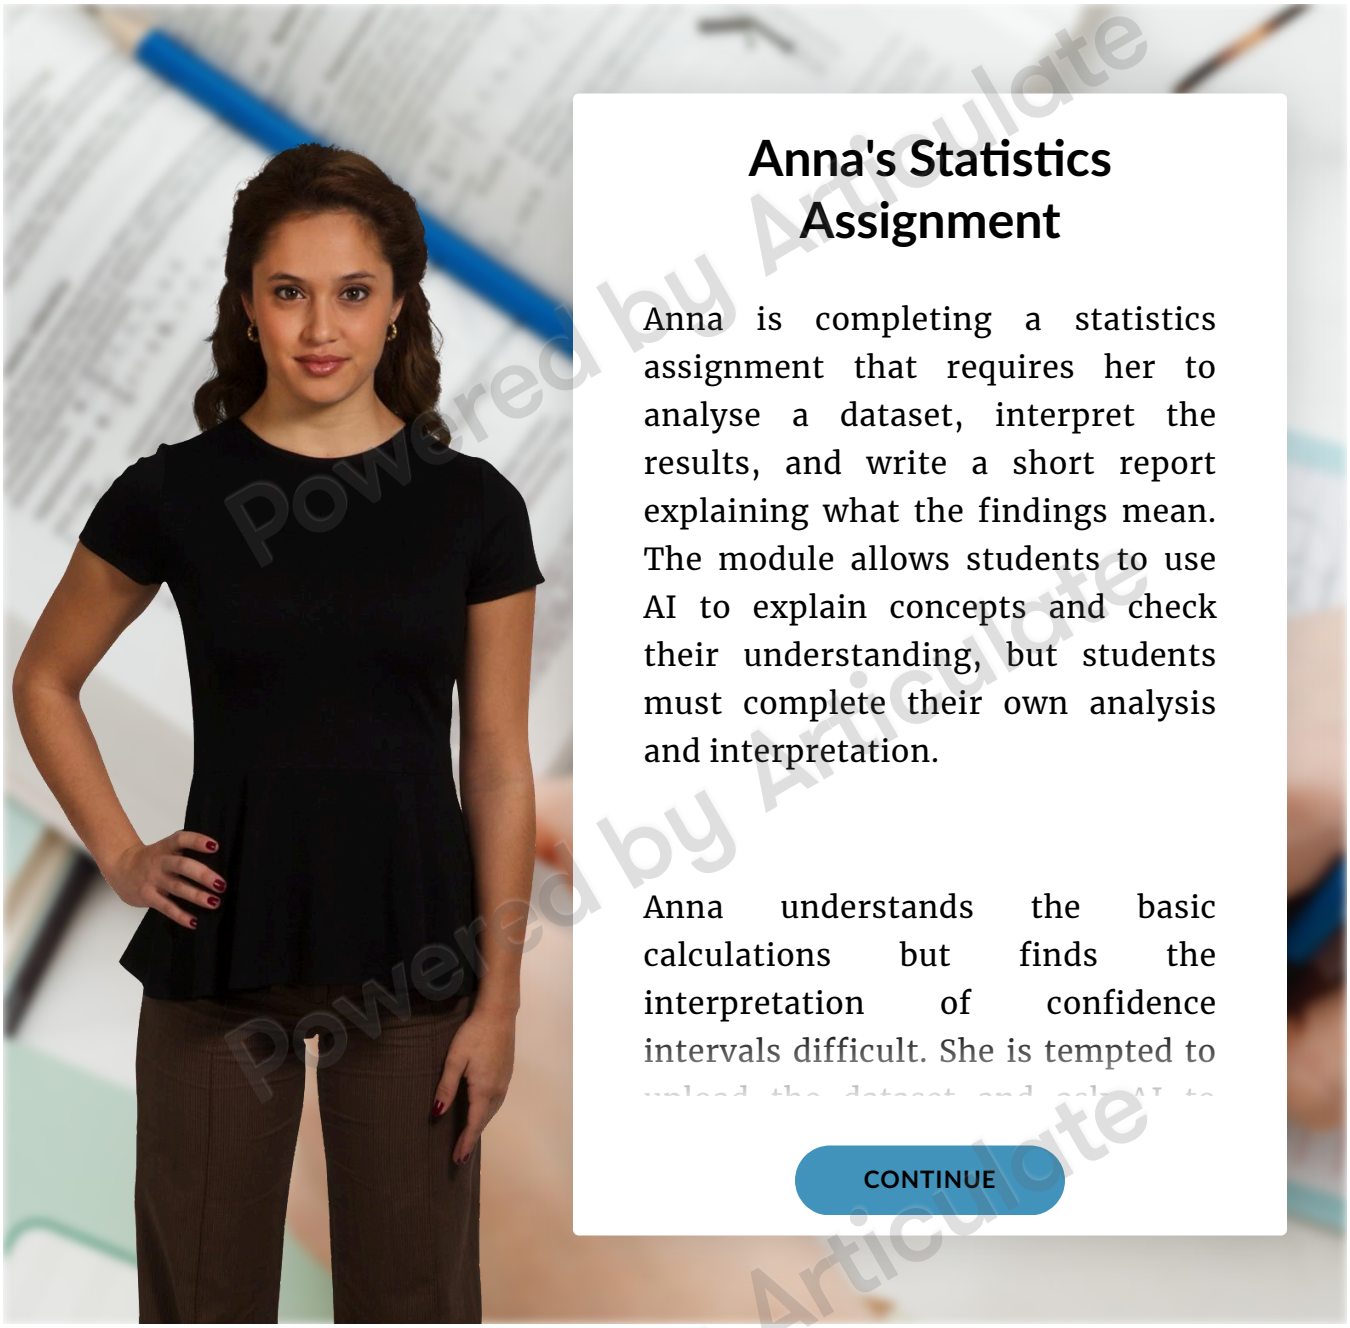

Powered by Articulate
Powered by Articulate
Powered by Articulate

Scene 2 Slide 1

Continue → Next Slide

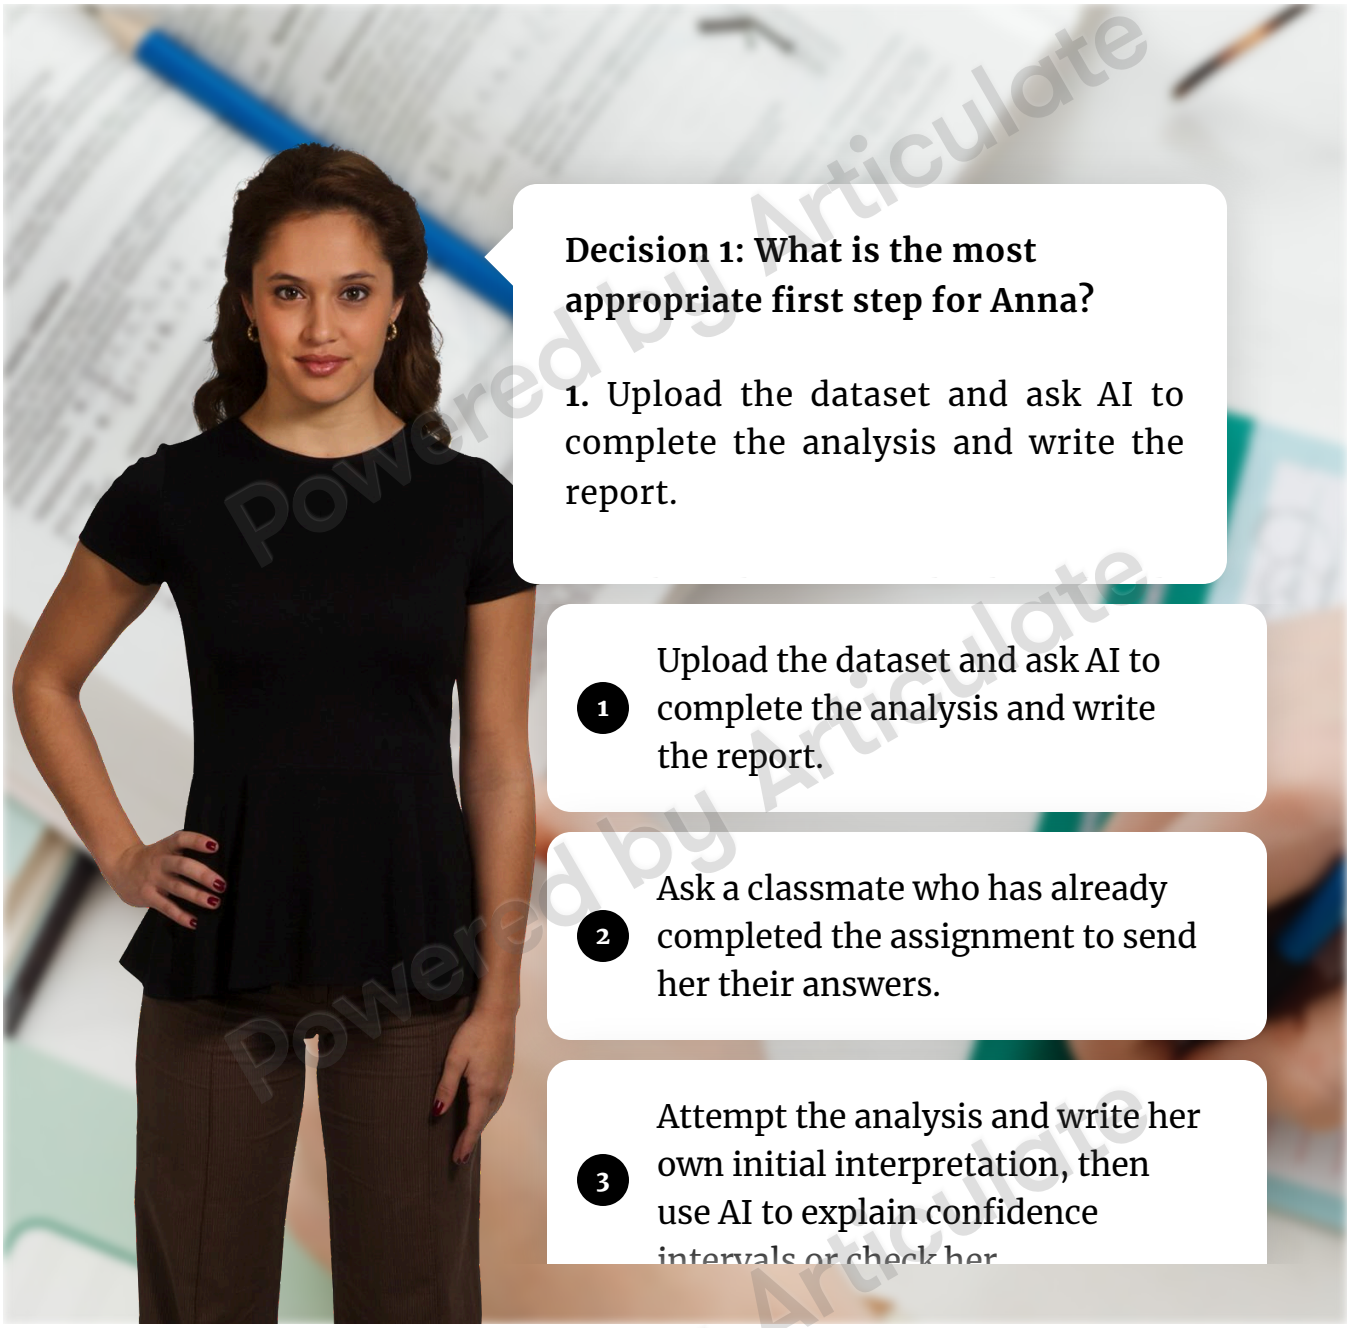

Powered by Articulate

Scene 2 Slide 2

0 → Next Slide

1 → Next Slide

2 → Next Slide

Powered by Articulate

Scene 2 Slide 3

0 → Next Slide

1 → Next Slide

2 → Next Slide

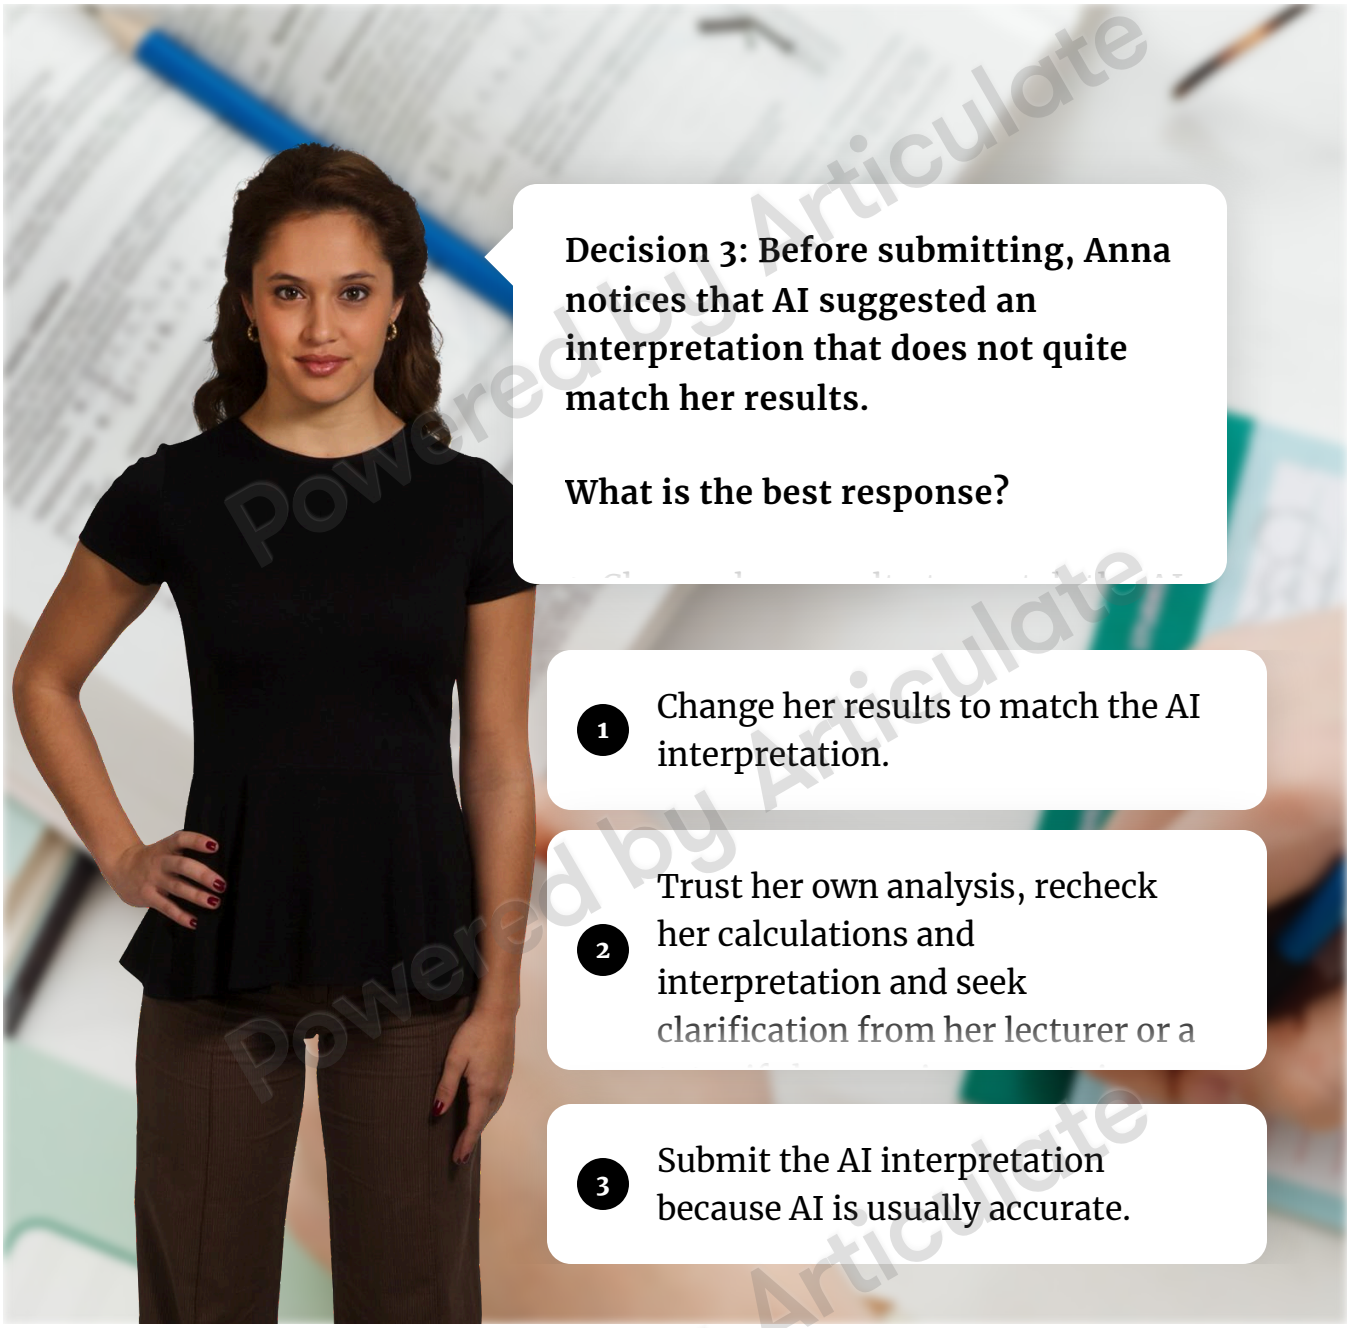

Powered by Articulate
Powered by Articulate
Powered by Articulate

Scene 2 Slide 4

0 → Next Slide

1 → Scene 1 Slide 5

2 → End of Scenario

Student Survey

NN Ntombana N Nhleko

Powered by Articulate

Well done on reaching the end of this course. Your feedback is very important to us and we would like to hear about your experience with this course, to help us evaluate your learning experience and improve the course.

How would you rate your overall experience of the course? *

	1	2	3	4	5	
Very poor	○	○	○	○	○	Excellent

How useful was the course in helping you understand how to use AI responsibly and ethically in higher education? *

Not at all	1	2	3	4	5	Very
	○	○	○	○	○	

How confident do you feel about applying what you learned about AI literacy to your studies? *

Not at all	1	2	3	4	5	Very
	○	○	○	○	○	

Is there anything else you would like to share about your experience of the course, or any suggestions for how it could be improved? *

0 / 300

SUBMIT

Powered by Articulate

Powered by Articulate
Powered by Articulate
Powered by Articulate